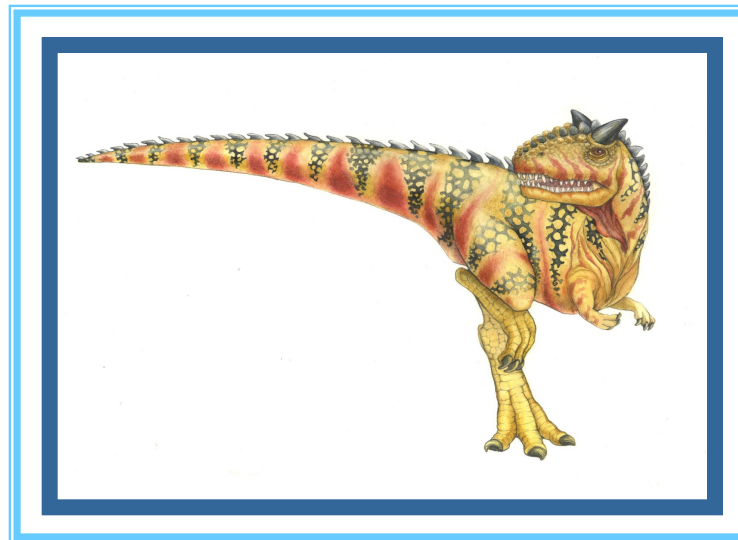


Chapter 7: Synchronization

Example





Synchronization Examples

- Classic Problems of Synchronization
 - Bounded-Buffer Problem
 - Readers and Writers Problem
- Window Synchronization
- POSIX Synchronization





Bounded-Buffer Problem

- n buffers, each can hold one item
- Semaphore `mutex` initialized to the value 1
- Semaphore `full` initialized to the value 0
- Semaphore `empty` initialized to the value n





Bounded Buffer Problem (Cont.)

- The structure of the producer process

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```





Bounded Buffer Problem (Cont.)

- The structure of the consumer process

```
do {  
    wait(full);  
    wait(mutex);  
  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex);  
    signal(empty);  
  
    ...  
    /* consume the item in next consumed */  
    ...  
} while (true);
```

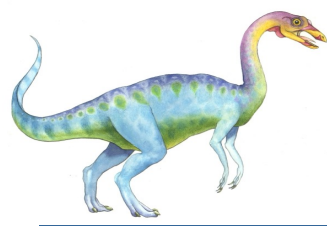




Readers-Writers Problem

- A data set is shared among a number of concurrent processes
 - **Readers** – only read the data; they **do not** perform any updates
 - **Writers** – can both read and write
- Problem – allow multiple readers to read the data set at the same time
 - Only one single writer can access shared data at a time
- Several variations of how readers and writers are treated – all involve different priorities.
- The **simplest** solution, referred to as the **first readers-writers problem**, requires that no reader be kept waiting unless a writer has already gained access to the shared data
 - Shared data update (by writers) can be delayed
 - This gives readers priority in accessing shared data
- Shared Data
 - Data set
 - Semaphore `rw_mutex` initialized to 1
 - Semaphore `mutex` initialized to 1
 - Integer `read_count` initialized to 0





Readers-Writers Problem (Cont.)

- The structure of a writer process

```
do {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
} while (true);
```





Readers-Writers Problem (Cont.)

■ The structure of a reader process

```
do {
    wait(mutex);
    read_count++;
    if (read_count == 1)
        wait(rw_mutex);
    signal(mutex)
    ...
    /* reading is performed */
    ...
    wait(mutex);
    read_count--;
    if (read_count == 0)
        signal(rw_mutex);
    signal(mutex);
} while (true);
```

Note:

- `rw_mutex` controls the access to shared data (critical section) for writers, and the first reader. The last reader leaving the critical section also has to release this lock
- `mutex` controls the access of readers to the shared variable `count`
- Writers wait on `rw_mutex`, first reader yet gain access to the critical section also waits on `rw_mutex`. All subsequent readers yet gain access wait on `mutex`





Readers-Writers Problem Variations

- **First variation** – no reader kept waiting unless a writer has gained access to use shared object. This is simple, but can result in starvation for writers, thus can potentially significantly delay the update of the object.
- **Second variation** – once a writer is ready, it needs to perform update asap. In another word, if a writer waits to access the object (this implies that there could be either readers or a writer inside), no new readers may start reading, i.e., they must wait after the writer updates the object
- A solution to either problem may result in starvation
- The problem can be solved or at least partially by kernel providing **reader-writer locks**, in which multiple processes are permitted to concurrently acquire a reader-writer lock in **read mode**, but only one process can acquire the reader-writer lock for writing (exclusive access). Acquiring a **reader-writer lock** thus requires specifying the mode of the lock: either **read** or **write** access





Synchronization Examples

- Solaris
- Windows XP
- Linux
- Pthreads





Solaris Synchronization


- Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing
- Uses **adaptive mutex** for efficiency when protecting data from *short code segments*, usually less than a few hundred (machine-level) instructions
 - Starts as a standard semaphore implemented as a **spinlock** in a multiprocessor system
 - If lock held, and by a thread running on another CPU, spins to wait for the lock to become available
 - If lock held by a non-run-state thread, block and sleep waiting for signal of lock being released
- Uses **condition variables**
- Uses **readers-writers locks** when longer sections of code need access to data. These are used to protect data that are frequently accessed, but usually in a read-only manner. The readers-writer locks are relatively expensive to implement.





Windows Synchronization

- The kernel uses interrupt masks to protect access to global resources in uniprocessor systems
- The kernel uses **spinlocks** in multiprocessor systems (to protect short code segments)
 - For efficiency, the kernel ensures that a thread will never be preempted while holding a spinlock
- For thread synchronization outside the kernel (user mode), Windows provides **dispatcher objects**, threads synchronize according to several different mechanisms, including mutex locks, semaphores, events, and timers
 - **Events** are similar to condition variables; they may notify a waiting thread when a desired condition occurs
 - **Timers** are used to notify one or more thread that a specified amount of time has expired
 - Dispatcher objects either **signaled-state** (object available) or **non-signaled state** (this means that another thread is holding the object, therefore the thread will block)





Linux Synchronization

- Linux:
 - Prior to kernel Version 2.6, disables interrupts to implement short critical sections
 - Version 2.6 and later, fully preemptive kernel
- Linux provides:
 - semaphores
 - Spinlocks – for multiprocessor systems
 - **atomic integer**, and all math operations using atomic integers performed without interruption
 - reader-writer locks
- On single-CPU system, spinlocks replaced by enabling and disabling kernel preemption



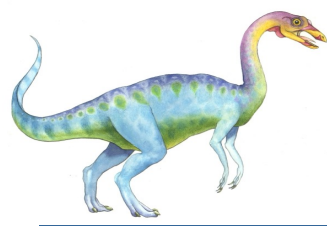


Atomic Variables

- **Atomic variables** - `atomic_t` is the type for atomic integer
- Consider the variables
`atomic_t counter;`
`int value;`

<i>Atomic Operation</i>	<i>Effect</i>
<code>atomic_set(&counter, 5);</code>	<code>counter = 5</code>
<code>atomic_add(10, &counter);</code>	<code>counter = counter + 10</code>
<code>atomic_sub(4, &counter);</code>	<code>counter = counter - 4</code>
<code>atomic_inc(&counter);</code>	<code>counter = counter + 1</code>
<code>value = atomic_read(&counter);</code>	<code>value = 12</code>





POSIX Synchronization

- POSIX API provides
 - mutex locks
 - semaphores
 - condition variables
- Widely used on UNIX, Linux, and MacOS





POSIX Mutex Locks

■ Creating and initializing the lock

```
#include <pthread.h>

pthread_mutex_t mutex;

/* create and initialize the mutex lock */
pthread_mutex_init(&mutex, NULL);
```

■ Acquiring and releasing the lock

```
/* acquire the mutex lock */
pthread_mutex_lock(&mutex);

/* critical section */

/* release the mutex lock */
pthread_mutex_unlock(&mutex);
```





POSIX Condition Variables

- POSIX condition variables are associated with a POSIX mutex lock to provide mutual exclusion: Creating and initializing the condition variable:

```
pthread_mutex_t mutex;  
pthread_cond_t cond_var;
```

```
pthread_mutex_init(&mutex, NULL);  
pthread_cond_init(&cond_var, NULL);
```





POSIX Condition Variables

- Thread waiting for the condition `a == b` to become true:

```
pthread_mutex_lock(&mutex);  
while (a != b)  
    pthread_cond_wait(&cond_var, &mutex);  
  
pthread_mutex_unlock(&mutex);
```

- Thread signaling another thread waiting on the condition variable:

```
pthread_mutex_lock(&mutex);  
a = b;  
pthread_cond_signal(&cond_var);  
pthread_mutex_unlock(&mutex);
```



End of Chapter 7

